

# Financial Fraud Analysis

## Import Libraries

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.metrics import classification_report, confusion_matrix, roc_auc_score
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from xgboost import XGBClassifier
```

## Load DataSet

```
In [2]: df = pd.read_csv("C:/Users/ankit/OneDrive/Desktop/Projects/Financial Fraud/financial
```

```
In [3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 10 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   transaction_id  10000 non-null   int64  
 1   user_id          10000 non-null   int64  
 2   amount           10000 non-null   float64 
 3   transaction_type 10000 non-null   object  
 4   merchant_category 10000 non-null   object  
 5   country          10000 non-null   object  
 6   hour              10000 non-null   int64  
 7   device_risk_score 10000 non-null   float64 
 8   ip_risk_score    10000 non-null   float64 
 9   is_fraud         10000 non-null   int64  
dtypes: float64(3), int64(4), object(3)
memory usage: 781.4+ KB
```

```
In [4]: df.head()
```

Out[4]:

	transaction_id	user_id	amount	transaction_type	merchant_category	country	I
0	9608	363	4922.587542	ATM	Travel	TR	
1	456	692	48.018303	QR	Food	US	
2	4747	587	136.881960	Online	Travel	TR	
3	6934	445	80.534719	POS	Clothing	TR	
4	1646	729	120.041158	Online	Grocery	FR	

In [5]: df.tail()

Out[5]:

	transaction_id	user_id	amount	transaction_type	merchant_category	country
9995	1076	482	58.366442	POS	Clothing	DE
9996	4995	904	139.502160	POS	Travel	DE
9997	3485	527	71.012122	Online	Travel	TR
9998	7922	771	21.031405	QR	Grocery	UK
9999	6451	429	54.028632	ATM	Electronics	DE

## Size of data

In [6]: df.dtypes

Out[6]:

transaction_id	int64
user_id	int64
amount	float64
transaction_type	object
merchant_category	object
country	object
hour	int64
device_risk_score	float64
ip_risk_score	float64
is_fraud	int64
dtype: object	

## Field Information

In [7]: df.columns

Out[7]:

```
Index(['transaction_id', 'user_id', 'amount', 'transaction_type',
       'merchant_category', 'country', 'hour', 'device_risk_score',
       'ip_risk_score', 'is_fraud'],
      dtype='object')
```

## Missing values and Duplicates

```
In [8]: print("Missing values per column:")
print(df.isnull().sum())

print("\nNumber of duplicate rows:", df.duplicated().sum())
```

Missing values per column:

	transaction_id	user_id	amount	transaction_type	merchant_category	country	hour	device_risk_score	ip_risk_score	is_fraud
count	0	0	0	0	0	0	0	0	0	0
std										
min										
25%										
50%										

dtype: int64

Number of duplicate rows: 0

```
In [9]: df.describe().T
```

```
Out[9]:
```

	count	mean	std	min	25%	50%
<b>transaction_id</b>	10000.0	4999.500000	2886.895680	0.000000	2499.750000	4999.500000
<b>user_id</b>	10000.0	500.058700	288.328495	0.000000	247.000000	503.000000
<b>amount</b>	10000.0	178.142763	531.647950	1.000000	65.084753	101.686510
<b>hour</b>	10000.0	14.247100	5.347383	0.000000	10.000000	14.000000
<b>device_risk_score</b>	10000.0	0.183773	0.177381	0.000030	0.075721	0.156583
<b>ip_risk_score</b>	10000.0	0.184669	0.175772	0.000009	0.077762	0.158290
<b>is_fraud</b>	10000.0	0.050000	0.217956	0.000000	0.000000	0.000000

```
In [10]: df.select_dtypes(include=np.number).corr()
```

```
Out[10]:
```

	transaction_id	user_id	amount	hour	device_risk_score	ip_risk_score	is_fraud
<b>transaction_id</b>	1.000000	-0.001360	0.242245	-0.070046	0.332211	0.332628	0.377492
<b>user_id</b>	-0.001360	1.000000	-0.008333	-0.022854	0.007316	0.004433	0.006165
<b>amount</b>	0.242245	-0.008333	1.000000	-0.100818	0.554977	0.549554	0.638435
<b>hour</b>	-0.070046	-0.022854	-0.100818	1.000000	-0.149832	-0.161670	-0.181448
<b>device_risk_score</b>	0.332211	0.007316	0.554977	-0.149832	1.000000	0.757978	0.871989
<b>ip_risk_score</b>	0.332628	0.004433	0.549554	-0.161670	0.757978	1.000000	0.871989
<b>is_fraud</b>	0.377492	0.006165	0.638435	-0.181448	0.871989	0.871989	1.000000

## Check the number of legitimate and fraudulent transactions and fraud ratio

```
In [11]: # Count of fraud vs Legitimate transactions
print(df['is_fraud'].value_counts())
```

```
# Fraud ratio
fraud_ratio = df['is_fraud'].mean()
print("Fraud ratio:", fraud_ratio)
```

```
is_fraud
0    9500
1     500
Name: count, dtype: int64
Fraud ratio: 0.05
```

## Basic statistics for numeric columns and compare fraud vs legitimate transactions

```
In [12]: # Describe numeric columns
print(df[['amount', 'hour', 'device_risk_score', 'ip_risk_score']].describe())
```

```
# Group by fraud
print(df.groupby('is_fraud')[['amount', 'hour']].mean())
```


	amount	hour	device_risk_score	ip_risk_score
count	10000.000000	10000.000000	10000.000000	10000.000000
mean	178.142763	14.247100	0.183773	0.184669
std	531.647950	5.347383	0.177381	0.175772
min	1.000000	0.000000	0.000030	0.000009
25%	65.084753	10.000000	0.075721	0.077762
50%	101.686510	14.000000	0.156583	0.158290
75%	138.280872	19.000000	0.234939	0.236968
max	11628.213881	23.000000	0.998737	0.999603
	amount	hour		
is_fraud				
0	100.277751	14.469684		
1	1657.577984	10.018000		

## Distribution of transaction amounts for fraud and legitimate transactions

```
In [13]: # Amount distribution by fraud
```

```
plt.figure(figsize=(10,5))
sns.histplot(data=df, x='amount', hue='is_fraud', bins=50, log_scale=True, palette='viridis')
plt.title("Transaction Amount Distribution by Fraud")
plt.xlabel("Amount")
plt.ylabel("Count")
plt.show()
```

```
C:\Users\ankit\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119: FutureWarning:
use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.
with pd.option_context('mode.use_inf_as_na', True):
```



## Number of transactions per hour, separated by fraud and legitimate labels

```
In [14]: df['is_fraud_label'] = df['is_fraud'].map({0: 'Non-Fraud', 1: 'Fraud'})


# Plot
plt.figure(figsize=(10, 5))

ax = sns.countplot(
    data=df,
    x='hour',
    hue='is_fraud_label',
    palette=['green', 'red']
)

# Titles and labels
plt.title("Hourly Transaction Distribution by Fraud Status", fontsize=14)
plt.xlabel("Hour of Day", fontsize=12)
plt.ylabel("Number of Transactions", fontsize=12)

ax.legend(title="Transaction Type")


plt.tight_layout()
plt.show()
```



## Fraud vs Non-Fraud Distribution Across Transaction Types and Merchant Categories


```
In [15]: grid = sns.FacetGrid(df, col='is_fraud', row='transaction_type', aspect=1.6)
grid.map(plt.hist, 'merchant_category', alpha=.5, bins=20)
grid.add_legend()
```

```
Out[15]: <seaborn.axisgrid.FacetGrid at 0x25188367550>
```



## Fraud Rate by Merchant Category


```
In [16]: fraud_by_cat = df.groupby('merchant_category')['is_fraud'].mean().sort_values(as
fraud_by_cat.plot.barh(figsize=(6,8))
plt.title('fraud rate by merchant_category')
plt.show()
```



## Distribution of Device and IP Risk Scores


```
In [17]: fig, ax = plt.subplots(1,2, figsize=(12,5))
sns.histplot(df['device_risk_score'], kde=True, ax=ax[0])
sns.histplot(df['ip_risk_score'], kde=True, ax=ax[1])
ax[0].set_title("Device Risk Score")
ax[1].set_title("IP Risk Score")
plt.show()
```

```
C:\Users\ankit\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119: FutureWarning:
use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
C:\Users\ankit\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119: FutureWarning:
use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
```



## Distribution of Transaction Types


```
In [18]: sns.countplot(data=df, x="transaction_type")
plt.title("Transaction Type Count")
plt.xticks(rotation=45)
plt.show()
```



## Merchant Category Distribution

```
In [19]: plt.figure(figsize=(10,6))
sns.countplot(data=df, y="merchant_category", order=df[ 'merchant_category' ].value
```


```
plt.title("Merchant Category Distribution")
plt.show()
```



## Fraud and legitimate transaction counts per country.

```
In [20]: # Pivot table: country vs fraud
country_fraud = df.pivot_table(index='country', columns='is_fraud', values='tran

# Plot heatmap
plt.figure(figsize=(8,5))
sns.heatmap(country_fraud, annot=True, fmt='g', cmap='coolwarm')
plt.title("Country vs Fraud Count")
plt.xlabel("Is Fraud")
plt.ylabel("Country")
plt.show()
```



## Data Modelling

```
In [21]: # Feature Engineering

def time_range(x):
    if 5 <= x <= 11:
        return "Morning"
    elif 12 <= x <= 17:
        return "Afternoon"
    elif 18 <= x <= 22:
        return "Evening"
    else:
        return "Night"

df["time_range"] = df["hour"].apply(time_range)

cat_cols = ["transaction_type", "merchant_category", "country", "time_range"]
le = LabelEncoder()
for col in cat_cols:
    df[col] = le.fit_transform(df[col])
```

### Prepare data

```
In [22]: X = df.drop(
    columns=["is_fraud", "is_fraud_label", "transaction_id"],
    errors="ignore")
y = df["is_fraud"]

#Encode categorical columns FIRST
```

```

cat_cols = X.select_dtypes(include="object").columns
le = LabelEncoder()

for col in cat_cols:
    X[col] = le.fit_transform(X[col])

feature_names = X.columns

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

X_train, X_test, y_train, y_test = train_test_split(
    X_scaled, y, test_size=0.25, random_state=42, stratify=y
)

```

## Train Multiple Model

In [23]:

```

lr = LogisticRegression(max_iter=1000)
lr.fit(X_train, y_train)

rf = RandomForestClassifier(random_state=42)
rf.fit(X_train, y_train)

dt = DecisionTreeClassifier(random_state=42)
dt.fit(X_train, y_train)

knn = KNeighborsClassifier()
knn.fit(X_train, y_train)

xgb = XGBClassifier(eval_metric="logloss", random_state=42)
xgb.fit(X_train, y_train)

```

Out[23]:

```

▼          XGBClassifier
XGBClassifier(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=None, device=None, early_stopping_rou
nds=None,
              enable_categorical=False, eval_metric='logloss',
              feature_types=None, gamma=None, grow_policy=None,
              importance_type=None, interaction_constraints=None,
              learning_rate=None, max_bin=None, max_cat_threshold>No
ne,

```

## Evaluate Models

In [24]:

```

def evaluate_model(model):
    pred = model.predict(X_test)
    print(classification_report(y_test, pred))
    cm = confusion_matrix(y_test, pred)

```

```
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```


## Evaluate all model

```
In [25]: models = [lr, rf, dt, knn, xgb]
names = ["Logistic Regression", "Random Forest", "Decision Tree", "KNN", "XGBoos"]

for name, model in zip(names, models):
    print(f"----- {name} -----")
    evaluate_model(model)
```


```
----- Logistic Regression -----
      precision    recall   f1-score   support
          0         1.00     1.00     1.00     2375
          1         1.00     1.00     1.00      125

      accuracy                           1.00     2500
      macro avg       1.00     1.00     1.00     2500
  weighted avg       1.00     1.00     1.00     2500
```




----- Random Forest -----

	precision	recall	f1-score	support
0	1.00	1.00	1.00	2375
1	1.00	1.00	1.00	125
accuracy			1.00	2500
macro avg	1.00	1.00	1.00	2500
weighted avg	1.00	1.00	1.00	2500




----- Decision Tree -----

	precision	recall	f1-score	support
0	1.00	1.00	1.00	2375
1	1.00	1.00	1.00	125
accuracy			1.00	2500
macro avg	1.00	1.00	1.00	2500
weighted avg	1.00	1.00	1.00	2500




----- KNN -----				
	precision	recall	f1-score	support
0	1.00	1.00	1.00	2375
1	1.00	1.00	1.00	125
accuracy			1.00	2500
macro avg	1.00	1.00	1.00	2500
weighted avg	1.00	1.00	1.00	2500



----- XGBoost -----

	precision	recall	f1-score	support
0	1.00	1.00	1.00	2375
1	1.00	1.00	1.00	125
accuracy			1.00	2500
macro avg	1.00	1.00	1.00	2500
weighted avg	1.00	1.00	1.00	2500




## ROC–AUC Curve (Best Model – XGBoost)

```
In [26]: y_pred_prob = xgb.predict_proba(X_test)[:, 1]
fpr, tpr, _ = roc_curve(y_test, y_pred_prob)

plt.plot(fpr, tpr, label="XGBoost")
plt.plot([0, 1], [0, 1], '--')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC-AUC Curve")
plt.legend()
plt.show()

roc_auc_score(y_test, y_pred_prob)
```



Out[26]: 1.0

## Predict fraud for 5 example transactions and visualize the results

```
In [28]: # Example new transactions
new_transactions = pd.DataFrame({
    'amount': [1200, 50, 5000, 300, 800],
    'hour': [2, 14, 23, 10, 1],
    'device_risk_score': [0.9, 0.1, 0.95, 0.2, 0.8],
    'ip_risk_score': [0.85, 0.05, 0.9, 0.15, 0.75]
})

# Create full feature frame using training features
new_X = pd.DataFrame(columns=X.columns)

# Fill known values
for col in new_transactions.columns:
    new_X[col] = new_transactions[col]

# Fill missing columns with 0
new_X = new_X.fillna(0)

new_X_scaled = scaler.transform(new_X)

# Make predictions
predictions = rf.predict(new_X_scaled)

new_transactions['predicted_is_fraud'] = predictions
new_transactions['predicted_label'] = new_transactions['predicted_is_fraud'].map
```

```

        {1: 'Fraud', 0: 'Legit'}
    )

print(new_transactions)

# Visualization


plt.figure(figsize=(10,6))
sns.barplot(
    x=new_transactions.index,
    y='amount',
    hue='predicted_label',
    data=new_transactions,
    dodge=False,
    palette=['green','red']
)

plt.title("Predicted Fraud vs Legit for Example Transactions")
plt.xlabel("Example Transaction")
plt.ylabel("Transaction Amount")
plt.show()

```

	amount	hour	device_risk_score	ip_risk_score	predicted_is_fraud	\
0	1200	2		0.90	0.85	1
1	50	14		0.10	0.05	0
2	5000	23		0.95	0.90	1
3	300	10		0.20	0.15	0
4	800	1		0.80	0.75	1


	predicted_label
0	Fraud
1	Legit
2	Fraud
3	Legit
4	Fraud



## Feature Importance (XGBoost)

```
In [29]: # Create feature importance DataFrame
importance_df = pd.DataFrame({
    "Feature": feature_names,
    "Importance": xgb.feature_importances_
}).sort_values(by="Importance", ascending=False)

plt.figure(figsize=(10,6))
sns.barplot(
    data=importance_df,
    x="Importance",
    y="Feature"
)
plt.title("Feature Importance (XGBoost)")
plt.xlabel("Importance Score")
plt.ylabel("Feature")
plt.tight_layout()
plt.show()
```



```
In [ ]:
```